

✓ Waste Classification - CNN Assignment

This notebook implements a comprehensive waste classification system including:

- **Part 1:** Custom CNN with optimizer comparisons (AdamW vs SGD vs SGD+Momentum)
- **Part 2:** Fine-tuned pretrained models (ResNet-18)
- **Part 3:** Performance comparison and analysis

Dataset: RealWaste - 9 waste categories

Instructions

How to run this notebook:

1. Execute cells sequentially from top to bottom
2. Training takes approximately 15-20 minutes per model (Adam, SGD, SGD+Momentum)
3. ResNet-18 fine-tuning takes additional 20-30 minutes
4. Total estimated time: 90-120 minutes on GPU

Key improvements implemented:

- Gentler data augmentation (rotation 15°, ColorJitter 0.2, scale 0.8-1.0)
- Class-weighted loss function for handling imbalance
- Dropout 0.4 for regularization
- ReduceLROnPlateau scheduler for smooth convergence

Expected results:

- Custom CNN: 65-75% test accuracy
- ResNet-18 Fine-tuned: 75-85% test accuracy

✓ Part 1: Custom CNN Architecture

Q1-Q2: Environment Setup and Dataset Loading

```
import json
import math
import random
from pathlib import Path
import time
```

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import torch
import torch.nn as nn
import torch.nn.functional as F
import torchvision.transforms as T
from torch.utils.data import DataLoader, Dataset
from torchvision import datasets
from sklearn.metrics import confusion_matrix, classification_report

DEVICE = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
SEED = 42
random.seed(SEED)
np.random.seed(SEED)
torch.manual_seed(SEED)
print(f'Using device: {DEVICE}')
```

```
Using device: cuda
```

```
!rm -rf work/
```

```
from pathlib import Path
import json
import random

!git clone --depth 1 --filter=blob:none https://github.com/Savidilis/realwaste

PROJECT_ROOT = Path('/content')
DATA_ROOT = PROJECT_ROOT / 'work' / 'Dataset' / 'RealWaste'
SPLIT_MANIFEST = PROJECT_ROOT / 'work' / 'Dataset' / 'realwaste_splits.json'
STATS_CACHE = PROJECT_ROOT / 'work' / 'Dataset' / 'realwaste_stats.json'
ARTIFACTS = PROJECT_ROOT / 'artifacts'
ARTIFACTS.mkdir(exist_ok=True)

print("DATA_ROOT =", DATA_ROOT)
assert DATA_ROOT.exists(), f'Dataset not found at {DATA_ROOT}'

# Create realwaste_splits.json
if not SPLIT_MANIFEST.exists():
    print("realwaste_splits.json not found – creating a new split manifest")
    # class names from folder names (stable alphabetical order)
    class_names = sorted([d.name for d in DATA_ROOT.iterdir() if d.is_dir()])
    random.seed(42)

    train_paths, val_paths, test_paths = [], [], []

    for cls_name in class_names:
        cls_dir = DATA_ROOT / cls_name
        # collect all files for this class
        cls_files = [
            f for f in cls_dir.glob('*.*') if f.is_file()
        ]
```

```
p.relative_to(DATA_ROOT)
    for p in cls_dir.glob('*')
        if p.is_file()
    ]
cls_files = [str(p) for p in cls_files]
cls_files.sort()
random.shuffle(cls_files)

n = len(cls_files)
n_train = int(0.7 * n)
n_val = int(0.15 * n)
# rest goes to test
train_paths.extend(cls_files[:n_train])
val_paths.extend(cls_files[n_train:n_train + n_val])
test_paths.extend(cls_files[n_train + n_val:])

manifest = {
    "class_names": class_names,
    "splits": {
        "train": train_paths,
        "val": val_paths,
        "test": test_paths,
    },
}

with SPLIT_MANIFEST.open("w") as f:
    json.dump(manifest, f, indent=2)
else:
    with SPLIT_MANIFEST.open() as f:
        manifest = json.load(f)

class_names = manifest['class_names']
print(f'Number of classes: {len(class_names)}')
print(f'Classes: {class_names}')
split_sizes = {k: len(v) for k, v in manifest['splits'].items()}
print(f'Train: {split_sizes["train"]}, Val: {split_sizes["val"]}, T
```

```
Cloning into 'work'...
remote: Enumerating objects: 13, done.
remote: Counting objects: 100% (13/13), done.
remote: Compressing objects: 100% (13/13), done.
remote: Total 13 (delta 0), reused 13 (delta 0), pack-reused 0 (from
Receiving objects: 100% (13/13), 114.33 KiB | 2.24 MiB/s, done.
remote: Enumerating objects: 4755, done.
remote: Counting objects: 100% (4755/4755), done.
remote: Compressing objects: 100% (4755/4755), done.
remote: Total 4755 (delta 0), reused 4755 (delta 0), pack-reused 0 (from
Receiving objects: 100% (4755/4755), 831.40 KiB | 7.99 MiB/s, done.
Updating files: 100% (4755/4755), done.
Filtering content: 100% (4752/4752), 656.42 MiB | 7.19 MiB/s, done.
DATA_ROOT = /content/work/Dataset/RealWaste
realwaste_splits.json not found – creating a new split manifest...
Number of classes: 9
Classes: ['Cardboard', 'Food Organics', 'Glass', 'Metal', 'Miscellar
Train: 3323, Val: 710, Test: 719
```

```
# load normalization statistics
if STATS_CACHE.exists():
    with STATS_CACHE.open() as f:
        stats = json.load(f)
        mean, std = stats['mean'], stats['std']
else:
    mean, std = [0.485, 0.456, 0.406], [0.229, 0.224, 0.225]

print(f'Normalization - Mean: {mean}, Std: {std}')
```

```
Normalization - Mean: [0.485, 0.456, 0.406], Std: [0.229, 0.224, 0.225]
```

▼ Q3: Data Augmentation and Transforms

```
# data augmentation and transforms
IMAGE_SIZE = 224
BATCH_SIZE = 32

train_transform = T.Compose([
    T.Resize((IMAGE_SIZE + 32, IMAGE_SIZE + 32)),
    T.RandomResizedCrop(IMAGE_SIZE, scale=(0.8, 1.0)),
    T.RandomHorizontalFlip(),
    T.RandomRotation(15),
    T.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2, hue=0.2),
    T.ToTensor(),
    T.Normalize(mean=mean, std=std),
])

eval_transform = T.Compose([
    T.Resize((IMAGE_SIZE, IMAGE_SIZE)),
    T.ToTensor(),
    T.Normalize(mean=mean, std=std)
])

print('Data augmentation configured')
```

```
Data augmentation configured
```

```
# dataset class
class WasteDataset(Dataset):
    def __init__(self, root, samples, class_to_idx, transform=None):
        self.root = Path(root)
        self.samples = list(samples)
        self.class_to_idx = class_to_idx
        self.transform = transform
        self.loader = datasets.folder.default_loader

    def __len__(self):
        return len(self.samples)
```

```
def __getitem__(self, idx):
    rel_path = self.samples[idx]
    label_name = Path(rel_path).parts[0]
    target = self.class_to_idx[label_name]
    img_path = self.root / rel_path
    img = self.loader(img_path)
    if self.transform:
        img = self.transform(img)
    return img, target

class_to_idx = {name: idx for idx, name in enumerate(class_names)}

train_dataset = WasteDataset(DATA_ROOT, manifest['splits']['train'])
val_dataset = WasteDataset(DATA_ROOT, manifest['splits']['val'], cl
test_dataset = WasteDataset(DATA_ROOT, manifest['splits']['test'],

train_loader = DataLoader(train_dataset, batch_size=BATCH_SIZE, shu
val_loader = DataLoader(val_dataset, batch_size=BATCH_SIZE, shuffle=
test_loader = DataLoader(test_dataset, batch_size=BATCH_SIZE, shuff

print(f'Dataloaders created - batch size: {BATCH_SIZE}')
```

Dataloaders created - batch size: 32

```
# visualize some training samples
def show_samples(loader, mean, std, n=8):
    imgs, labels = next(iter(loader))
    inv_normalize = T.Normalize(
        mean=[-m/s for m, s in zip(mean, std)],
        std=[1/s for s in std]
    )

    imgs = imgs[:n]
    imgs = [inv_normalize(img).permute(1, 2, 0).clamp(0, 1).numpy()
    labs = [class_names[l.item()] for l in labels[:n]]

    cols = 4
    rows = math.ceil(len(imgs) / cols)


    fig, axes = plt.subplots(rows, cols, figsize=(12, 3*rows))
    axes = axes.flatten() if rows > 1 else [axes] if cols == 1 else

    for i, (img, lab) in enumerate(zip(imgs, labs)):
        axes[i].imshow(img)
        axes[i].set_title(lab)
        axes[i].axis('off')

    for i in range(len(imgs), len(axes)):
        axes[i].axis('off')

    plt.tight_layout()
    plt.show()

show_samples(train_loader, mean, std, n=8)
```



▼ Q4: Visualize Training Samples

```
# model definition - no batch normalization
class SimpleCNN(nn.Module):
    def __init__(self, num_classes=len(class_names)):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 32, kernel_size=3, padding=1)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, padding=1)
        self.conv3 = nn.Conv2d(64, 128, kernel_size=3, padding=1)
        self.pool = nn.MaxPool2d(2, 2)
        self.dropout = nn.Dropout(0.4)

        # calculate feature size after 3 pooling operations
        feature_size = 128 * (IMAGE_SIZE // 8) * (IMAGE_SIZE // 8)
        self.fc1 = nn.Linear(feature_size, 256)
        self.fc2 = nn.Linear(256, num_classes)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
```

```
        x = self.pool(F.relu(self.conv2(x)))
        x = self.pool(F.relu(self.conv3(x)))
        x = torch.flatten(x, 1)
        x = F.relu(self.fc1(x))
        x = self.dropout(x)
        x = self.fc2(x)
        return x

model = SimpleCNN().to(DEVICE)
total_params = sum(p.numel() for p in model.parameters())
print(f'Model created with {total_params:,} parameters')
print(model)

Model created with 25,785,929 parameters
SimpleCNN(
  (conv1): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=
  (conv2): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=
  (conv3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
  (dropout): Dropout(p=0.4, inplace=False)
  (fc1): Linear(in_features=100352, out_features=256, bias=True)
  (fc2): Linear(in_features=256, out_features=9, bias=True)
)
```

▼ Q5-Q6: Define Custom CNN Architecture

The model follows a simple pattern: convolution + ReLU + max-pool, repeated, then fully connected layers with dropout. **No batch normalization layers** are used as per assignment requirements.

Architecture choices:

- ReLU activation
- Kernel sizes: 3×3
- Filters: 32, 64, 128
- FC layer: 256 units
- Dropout: 0.4
- Class-weighted loss for handling imbalance

```
# training function
def train_model(model, train_loader, val_loader, epochs=20, lr=1e-3
    # compute class weights for balanced loss
    from collections import Counter
    train_labels = [Path(p).parts[0] for p in manifest['splits'][t]
    label_counts = Counter(train_labels)
    class_counts = torch.tensor([label_counts[name] for name in cla
    class_weights = 1.0 / class_counts
    class_weights = class_weights / class_weights.sum() * len(class
    class_weights = class_weights.to(DEVICE)

    ...
```

```
optimizer = torch.optim.Adam(model.parameters(), lr=lr, weight_
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer)
criterion = nn.CrossEntropyLoss(weight=class_weights)

history = []
best_val_acc = 0.0
best_state = None

for epoch in range(1, epochs + 1):
    start_time = time.time()

    # training phase
    model.train()
    train_loss = 0.0
    train_correct = 0
    train_total = 0

    for inputs, targets in train_loader:
        inputs, targets = inputs.to(DEVICE), targets.to(DEVICE)

        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, targets)
        loss.backward()
        optimizer.step()

        train_loss += loss.item() * inputs.size(0)
        _, predicted = outputs.max(1)
        train_correct += predicted.eq(targets).sum().item()
        train_total += targets.size(0)

    train_loss = train_loss / train_total
    train_acc = train_correct / train_total

    # validation phase
    model.eval()
    val_loss = 0.0
    val_correct = 0
    val_total = 0

    with torch.no_grad():
        for inputs, targets in val_loader:
            inputs, targets = inputs.to(DEVICE), targets.to(DEVICE)
            outputs = model(inputs)
            loss = criterion(outputs, targets)

            val_loss += loss.item() * inputs.size(0)
            _, predicted = outputs.max(1)
            val_correct += predicted.eq(targets).sum().item()
            val_total += targets.size(0)

    val_loss = val_loss / val_total
    val_acc = val_correct / val_total
```

.....

```
scneuler.step(val_loss)

history.append({
    'epoch': epoch,
    'train_loss': train_loss,
    'train_acc': train_acc,
    'val_loss': val_loss,
    'val_acc': val_acc
})

if val_acc > best_val_acc:
    best_val_acc = val_acc
    best_state = model.state_dict()

elapsed = time.time() - start_time
print(f'Epoch {epoch:02d}/{epochs} | '
      f'Train Loss: {train_loss:.4f} Acc: {train_acc:.3f} | '
      f'Val Loss: {val_loss:.4f} Acc: {val_acc:.3f} | '
      f'Time: {elapsed:.1f}s')

if best_state:
    model.load_state_dict(best_state)
    print(f'\nLoaded best model with validation accuracy: {best_val_acc:.3f}')

return history
```

▼ Q7: Training Function (AdamW Optimizer)

```
# evaluation function
def evaluate(model, loader):
    model.eval()
    all_targets = []
    all_predictions = []

    with torch.no_grad():
        for inputs, targets in loader:
            inputs = inputs.to(DEVICE)
            outputs = model(inputs)
            _, predicted = outputs.max(1)

            all_targets.extend(targets.cpu().numpy())
            all_predictions.extend(predicted.cpu().numpy())

    all_targets = np.array(all_targets)
    all_predictions = np.array(all_predictions)

    cm = confusion_matrix(all_targets, all_predictions)
    precision, recall, f1, _ = precision_recall_fscore_support(
        all_targets, all_predictions, average='macro', zero_division=0)
    accuracy = (all_predictions == all_targets).mean()
```

```
        return {
            'accuracy': float(accuracy),
            'precision': float(precision),
            'recall': float(recall),
            'f1': float(f1),
            'confusion_matrix': cm
        }
```

▼ Evaluation Function

```
# train the model
history = train_model(model, train_loader, val_loader, epochs=20, l

Epoch 01/20 | Train Loss: 2.1572 Acc: 0.218 | Val Loss: 1.7780 Acc:
Epoch 02/20 | Train Loss: 1.6971 Acc: 0.359 | Val Loss: 1.6126 Acc:
Epoch 03/20 | Train Loss: 1.5121 Acc: 0.423 | Val Loss: 1.3360 Acc:
Epoch 04/20 | Train Loss: 1.3727 Acc: 0.484 | Val Loss: 1.4467 Acc:
Epoch 05/20 | Train Loss: 1.3377 Acc: 0.485 | Val Loss: 1.4350 Acc:
Epoch 06/20 | Train Loss: 1.2909 Acc: 0.512 | Val Loss: 1.2334 Acc:
Epoch 07/20 | Train Loss: 1.2343 Acc: 0.527 | Val Loss: 1.2397 Acc:
Epoch 08/20 | Train Loss: 1.1897 Acc: 0.543 | Val Loss: 1.2765 Acc:
Epoch 09/20 | Train Loss: 1.2067 Acc: 0.543 | Val Loss: 1.0910 Acc:
Epoch 10/20 | Train Loss: 1.1455 Acc: 0.565 | Val Loss: 1.1237 Acc:
Epoch 11/20 | Train Loss: 1.1006 Acc: 0.569 | Val Loss: 1.1247 Acc:
Epoch 12/20 | Train Loss: 1.0850 Acc: 0.591 | Val Loss: 1.0525 Acc:
Epoch 13/20 | Train Loss: 1.0511 Acc: 0.600 | Val Loss: 1.0535 Acc:
Epoch 14/20 | Train Loss: 1.0162 Acc: 0.615 | Val Loss: 1.1347 Acc:
Epoch 15/20 | Train Loss: 0.9969 Acc: 0.612 | Val Loss: 1.0211 Acc:
Epoch 16/20 | Train Loss: 1.0008 Acc: 0.626 | Val Loss: 0.9259 Acc:
Epoch 17/20 | Train Loss: 0.9407 Acc: 0.638 | Val Loss: 0.9575 Acc:
Epoch 18/20 | Train Loss: 0.9498 Acc: 0.629 | Val Loss: 0.9724 Acc:
Epoch 19/20 | Train Loss: 0.9215 Acc: 0.636 | Val Loss: 0.8776 Acc:
Epoch 20/20 | Train Loss: 0.8786 Acc: 0.658 | Val Loss: 0.9406 Acc:

Loaded best model with validation accuracy: 0.686
```

▼ Q8: Train Model with Adam Optimizer (20 epochs)

Optimizer Used: Adam

- Adam adapts learning rates per parameter for faster convergence
- Works well with small datasets and noisy gradients
- More stable training compared to plain SGD

Class Weighting:

- Uses inverse class frequency weights in the loss function
- Helps the model pay more attention to underrepresented classes (Metal, Miscellaneous Trash)
- More stable than sample reweighting strategies


```
# plot training history
epochs = [h['epoch'] for h in history]
train_losses = [h['train_loss'] for h in history]
val_losses = [h['val_loss'] for h in history]
train_accs = [h['train_acc'] for h in history]
val_accs = [h['val_acc'] for h in history]

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 5))

ax1.plot(epochs, train_losses, 'b-', label='Train Loss')
ax1.plot(epochs, val_losses, 'r-', label='Val Loss')
ax1.set_xlabel('Epoch')
ax1.set_ylabel('Loss')
ax1.set_title('Training and Validation Loss')
ax1.legend()
ax1.grid(True, alpha=0.3)

ax2.plot(epochs, train_accs, 'b-', label='Train Acc')
ax2.plot(epochs, val_accs, 'r-', label='Val Acc')
ax2.set_xlabel('Epoch')
ax2.set_ylabel('Accuracy')
ax2.set_title('Training and Validation Accuracy')
ax2.legend()
ax2.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```



▼ Q9: Plot Training History

Learning Rate: 1e-3 with ReduceLROnPlateau scheduler

- Learning rate reduces by factor of 0.5 when validation loss plateaus for 3 epochs
- This adaptive approach prevents overfitting in later epochs
- Allows the model to start with faster learning and fine-tune with smaller steps

```
# evaluate on test set
print('Evaluating on test set...\n')
test_metrics = evaluate(model, test_loader)


print(f'Test Accuracy: {test_metrics["accuracy"]:.4f}')
print(f'Precision (macro): {test_metrics["precision"]:.4f}')
print(f'Recall (macro): {test_metrics["recall"]:.4f}')
print(f'F1 Score (macro): {test_metrics["f1"]:.4f}')


Evaluating on test set...

Test Accuracy: 0.6579
Precision (macro): 0.6541
Recall (macro): 0.6718
F1 Score (macro): 0.6573
```

▼ Q10: Test Set Evaluation (AdamW Model)

```
# confusion matrix
plt.figure(figsize=(12, 10))
sns.heatmap(
    test_metrics['confusion_matrix'],
    annot=True,
    fmt='d',
    cmap='Blues',
    xticklabels=class_names,
    yticklabels=class_names
)
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix - Test Set')
plt.xticks(rotation=45, ha='right')
plt.yticks(rotation=0)
plt.tight_layout()
plt.show()
```





✓ Confusion Matrix (AdamW Model)

```
# save model
```

```
model_path = ARTIFACTS / 'cnn_model.pt'
torch.save({
    'model_state_dict': model.state_dict(),
    'class_names': class_names,
    'test_accuracy': test_metrics['accuracy']
}, model_path)
print(f'Model saved to {model_path}')
```

```
Model saved to /content/artifacts/cnn_model.pt
```

▼ Q11-Q12: Compare Different Optimizers

Now we compare Adam against SGD and SGD with Momentum to analyze their performance differences.

```
# optimizer comparison - train with SGD and SGD+Momentum

# store Adam results
adam_history = {
    'train_loss': [h['train_loss'] for h in history],
    'val_loss': [h['val_loss'] for h in history],
    'train_acc': [h['train_acc'] for h in history],
    'val_acc': [h['val_acc'] for h in history]
}

def train_with_optimizer(optimizer_name, lr=0.01, momentum=0.0):
    """Train a new model instance with specified optimizer"""
    print(f'\n{"*50}')
    print(f'Training with {optimizer_name}')
    print(f'{{"*50}}')

    # create fresh model
    model_new = SimpleCNN().to(DEVICE)

    # compute class weights
    from collections import Counter
    train_labels = [Path(p).parts[0] for p in manifest['splits'][t]
    label_counts = Counter(train_labels)
    class_counts = torch.tensor([label_counts[name] for name in cla
    class_weights = 1.0 / class_counts
    class_weights = class_weights / class_weights.sum() * len(class
    class_weights = class_weights.to(DEVICE)

    # setup optimizer
    if optimizer_name == 'SGD':
        optimizer = torch.optim.SGD(model_new.parameters(), lr=lr,
    elif optimizer_name == 'SGD+Momentum':
        optimizer = torch.optim.SGD(model_new.parameters(), lr=lr,
                                    weight_decay=1e-4, nesterov=True
    else:
        raise ValueError(f"Unknown optimizer: {optimizer_name}")
```

```
scheduler = torch.optim.lr_scheduler.MultiStepLR(optimizer, mil
criterion = nn.CrossEntropyLoss(weight=class_weights)

history = []
best_val_acc = 0.0
best_state = None

for epoch in range(1, 21):
    start_time = time.time()

    # training phase
    model_new.train()
    train_loss = 0.0
    train_correct = 0
    train_total = 0

    for inputs, targets in train_loader:
        inputs, targets = inputs.to(DEVICE), targets.to(DEVICE)

        optimizer.zero_grad()
        outputs = model_new(inputs)
        loss = criterion(outputs, targets)
        loss.backward()
        nn.utils.clip_grad_norm_(model_new.parameters(), max_nc
optimizer.step()

        train_loss += loss.item() * inputs.size(0)
        _, predicted = outputs.max(1)
        train_correct += predicted.eq(targets).sum().item()
        train_total += targets.size(0)

    train_loss = train_loss / train_total
    train_acc = train_correct / train_total

    # validation phase
    model_new.eval()
    val_loss = 0.0
    val_correct = 0
    val_total = 0

    with torch.no_grad():
        for inputs, targets in val_loader:
            inputs, targets = inputs.to(DEVICE), targets.to(DEV
            outputs = model_new(inputs)
            loss = criterion(outputs, targets)

            val_loss += loss.item() * inputs.size(0)
            _, predicted = outputs.max(1)
            val_correct += predicted.eq(targets).sum().item()
            val_total += targets.size(0)

    val_loss = val_loss / val_total
    val_acc = val_correct / val_total
```

```
        scheduler.step()

        history.append({
            'epoch': epoch,
            'train_loss': train_loss,
            'train_acc': train_acc,
            'val_loss': val_loss,
            'val_acc': val_acc
        })

        if val_acc > best_val_acc:
            best_val_acc = val_acc
            best_state = model_new.state_dict()

        elapsed = time.time() - start_time
        print(f'Epoch {epoch:02d}/20 | '
              f'Train Loss: {train_loss:.4f} Acc: {train_acc:.3f} | '
              f'Val Loss: {val_loss:.4f} Acc: {val_acc:.3f} | '
              f'Time: {elapsed:.1f}s')

    if best_state:
        model_new.load_state_dict(best_state)
        print(f'\nLoaded best model with validation accuracy: {best_val_acc:.3f}')

    return model_new, history

# train with SGD
sgd_model, sgd_history = train_with_optimizer('SGD', lr=0.1, momentum=0)

# train with SGD+Momentum
sgdm_model, sgdm_history = train_with_optimizer('SGD+Momentum', lr=0.1, momentum=0.9)
```

```
=====
Training with SGD
=====

Epoch 01/20 | Train Loss: 2.0784 Acc: 0.198 | Val Loss: 1.8539 Acc: 0.198
Epoch 02/20 | Train Loss: 1.8142 Acc: 0.310 | Val Loss: 1.6632 Acc: 0.310
Epoch 03/20 | Train Loss: 1.6730 Acc: 0.377 | Val Loss: 1.6416 Acc: 0.377
Epoch 04/20 | Train Loss: 1.5849 Acc: 0.389 | Val Loss: 1.6039 Acc: 0.389
Epoch 05/20 | Train Loss: 1.5156 Acc: 0.438 | Val Loss: 1.4262 Acc: 0.438
Epoch 06/20 | Train Loss: 1.4633 Acc: 0.450 | Val Loss: 1.5322 Acc: 0.450
Epoch 07/20 | Train Loss: 1.4315 Acc: 0.468 | Val Loss: 1.4292 Acc: 0.468
Epoch 08/20 | Train Loss: 1.3796 Acc: 0.489 | Val Loss: 1.3726 Acc: 0.489
Epoch 09/20 | Train Loss: 1.3498 Acc: 0.497 | Val Loss: 1.3331 Acc: 0.497
Epoch 10/20 | Train Loss: 1.2945 Acc: 0.507 | Val Loss: 1.3725 Acc: 0.507
Epoch 11/20 | Train Loss: 1.2062 Acc: 0.547 | Val Loss: 1.2168 Acc: 0.547
Epoch 12/20 | Train Loss: 1.1819 Acc: 0.556 | Val Loss: 1.2005 Acc: 0.556
Epoch 13/20 | Train Loss: 1.1459 Acc: 0.563 | Val Loss: 1.2238 Acc: 0.563
Epoch 14/20 | Train Loss: 1.1497 Acc: 0.556 | Val Loss: 1.1974 Acc: 0.556
Epoch 15/20 | Train Loss: 1.1453 Acc: 0.570 | Val Loss: 1.1896 Acc: 0.570
Epoch 16/20 | Train Loss: 1.1350 Acc: 0.577 | Val Loss: 1.1851 Acc: 0.577
Epoch 17/20 | Train Loss: 1.1460 Acc: 0.562 | Val Loss: 1.1932 Acc: 0.562
Epoch 18/20 | Train Loss: 1.1301 Acc: 0.574 | Val Loss: 1.2005 Acc: 0.574
Epoch 19/20 | Train Loss: 1.1314 Acc: 0.569 | Val Loss: 1.1985 Acc: 0.569
```

```
Epoch 20/20 | Train Loss: 1.1103 ACC: 0.580 | Val Loss: 1.1991 ACC:
```

```
Loaded best model with validation accuracy: 0.606
```

```
=====
```

```
Training with SGD+Momentum
```

```
=====
```

Epoch 01/20	Train Loss: 2.1430 Acc: 0.188	Val Loss: 1.9095 Acc:
Epoch 02/20	Train Loss: 2.0511 Acc: 0.233	Val Loss: 1.9086 Acc:
Epoch 03/20	Train Loss: 1.9402 Acc: 0.286	Val Loss: 1.7798 Acc:
Epoch 04/20	Train Loss: 1.7883 Acc: 0.336	Val Loss: 1.7379 Acc:
Epoch 05/20	Train Loss: 1.7289 Acc: 0.352	Val Loss: 1.7614 Acc:
Epoch 06/20	Train Loss: 1.6764 Acc: 0.359	Val Loss: 1.6294 Acc:
Epoch 07/20	Train Loss: 1.6076 Acc: 0.396	Val Loss: 1.8067 Acc:
Epoch 08/20	Train Loss: 1.5490 Acc: 0.401	Val Loss: 1.6113 Acc:
Epoch 09/20	Train Loss: 1.5462 Acc: 0.414	Val Loss: 1.5447 Acc:
Epoch 10/20	Train Loss: 1.5093 Acc: 0.421	Val Loss: 1.6124 Acc:
Epoch 11/20	Train Loss: 1.3236 Acc: 0.471	Val Loss: 1.3791 Acc:
Epoch 12/20	Train Loss: 1.2041 Acc: 0.524	Val Loss: 1.4063 Acc:
Epoch 13/20	Train Loss: 1.1982 Acc: 0.548	Val Loss: 1.2828 Acc:
Epoch 14/20	Train Loss: 1.1743 Acc: 0.553	Val Loss: 1.2392 Acc:
Epoch 15/20	Train Loss: 1.1398 Acc: 0.560	Val Loss: 1.2229 Acc:
Epoch 16/20	Train Loss: 1.1256 Acc: 0.559	Val Loss: 1.2089 Acc:
Epoch 17/20	Train Loss: 1.1132 Acc: 0.566	Val Loss: 1.2122 Acc:
Epoch 18/20	Train Loss: 1.1028 Acc: 0.578	Val Loss: 1.2161 Acc:
Epoch 19/20	Train Loss: 1.1112 Acc: 0.573	Val Loss: 1.2031 Acc:
Epoch 20/20	Train Loss: 1.0817 Acc: 0.579	Val Loss: 1.2092 Acc:

```
Loaded best model with validation accuracy: 0.586
```

```
# plot optimizer comparison
import pandas as pd

fig, axes = plt.subplots(1, 2, figsize=(14, 5))
epochs_range = range(1, 21)

# validation accuracy
axes[0].plot(epochs_range, adam_history['val_acc'], 'b-', label='Adam')
axes[0].plot(epochs_range, [h['val_acc'] for h in sgd_history], 'r-')
axes[0].plot(epochs_range, [h['val_acc'] for h in sgdm_history], 'g-')
axes[0].set_xlabel('Epoch')
axes[0].set_ylabel('Validation Accuracy')
axes[0].set_title('Validation Accuracy Comparison')
axes[0].legend()
axes[0].grid(True, alpha=0.3)

# validation loss
axes[1].plot(epochs_range, adam_history['val_loss'], 'b-', label='Adam')
axes[1].plot(epochs_range, [h['val_loss'] for h in sgd_history], 'r-')
axes[1].plot(epochs_range, [h['val_loss'] for h in sgdm_history], 'g-')
axes[1].set_xlabel('Epoch')
axes[1].set_ylabel('Validation Loss')
axes[1].set_title('Validation Loss Comparison')
axes[1].legend()
axes[1].grid(True, alpha=0.3)
```

```
.....
```


```

plt.tight_layout()
plt.show()

# summary table
summary_data = {
    'Optimizer': ['Adam', 'SGD', 'SGD+Momentum'],
    'Best Val Acc': [
        max(adam_history['val_acc']),
        max([h['val_acc'] for h in sgd_history]),
        max([h['val_acc'] for h in sgdm_history])
    ],
    'Best Val Loss': [
        min(adam_history['val_loss']),
        min([h['val_loss'] for h in sgd_history]),
        min([h['val_loss'] for h in sgdm_history])
    ]
}
}

summary_df = pd.DataFrame(summary_data).sort_values('Best Val Acc', . . .
print('\n' + '='*60)
print('Optimizer Comparison Summary')
print('='*60)
print(summary_df.to_string(index=False))
print('='*60)

```



Optimizer Comparison Summary

Optimizer	Best Val Acc	Best Val Loss
Adam	0.685915	0.877583
SGD	0.605634	1.185129
SGD+Momentum	0.585915	1.203120

Q13: Analysis of Optimizer Results

SGD (no momentum):

- Steps strictly follow the current batch gradient
- Learning is slow and noisy
- Often gets stuck in local minima
- Convergence is slower compared to other optimizers

SGD+Momentum ($\beta=0.9$):

- Keeps a running average of past updates
- Smooths out zig-zags in the optimization path
- Accelerates progress along stable directions
- Usually reaches higher validation accuracy than plain SGD
- Better at escaping shallow local minima

Adam:

- Adapts the step size per parameter (adaptive learning rates)
- Combines benefits of momentum and RMSprop
- Best validation performance in our experiments
- More stable training with faster convergence
- Works well with small datasets and noisy gradients

Save AdamW Model

▼ Part 2: Fine-tuning Pretrained Models

Q14: Select and Load Pretrained Model

We will use **ResNet-18** pretrained on ImageNet. ResNet-18 is a proven architecture with:

- 18 layers deep with residual connections
- Trained on 1000 ImageNet classes
- Strong feature extraction capabilities
- Good balance between performance and computational cost

```
# load pretrained ResNet-18
from torchvision import models

# load with ImageNet pretrained weights
resnet18 = models.resnet18(weights=models.ResNet18_Weights.DEFAULT)

# replace final fully connected layer for our dataset
num_features = resnet18.fc.in_features
resnet18.fc = nn.Sequential(
    nn.Dropout(p=0.35),
    nn.Linear(num_features, len(class_names))
)

resnet18 = resnet18.to(DEVICE)

# count parameters
total_params = sum(p.numel() for p in resnet18.parameters())
trainable_params = sum(p.numel() for p in resnet18.parameters() if

print(f'ResNet-18 loaded with ImageNet weights')
print(f'Total parameters: {total_params:,}')
print(f'Trainable parameters: {trainable_params:,}')
print(f'Output classes: {len(class_names)}')

Downloading: "https://download.pytorch.org/models/resnet18-f37072fd.100%|\[REDACTED\]| 44.7M/44.7M \[00:00<00:00, 162MB/s\]
ResNet-18 loaded with ImageNet weights
Total parameters: 11,181,129
Trainable parameters: 11,181,129
Output classes: 9
```

▼ Q15-Q16: Fine-tuning Strategy and Training

Fine-tuning Approach:

1. **Phase A (Warm-up):** Freeze backbone, train only the new classification head
 - Allows new head to learn without disrupting pretrained features
 - 8 epochs with higher learning rate (6e-4)
2. **Phase B (Full fine-tuning):** Unfreeze all layers with discriminative learning rates
 - Backbone: lower LR (5e-5) to preserve pretrained features
 - Head: higher LR (1e-4) for faster adaptation
 - Cosine annealing scheduler for smooth convergence

Regularization techniques:

- Dropout (0.35) in classification head

- Label smoothing (0.03) to prevent overconfidence
- Gradient clipping (max_norm=1.0)
- Weight decay (2e-4 backbone, 6e-4 head)
- Batch normalization layers frozen in eval mode

```
# fine-tuning functions and utilities

def freeze_bn_stats(model):
    """Freeze batch normalization layers"""
    for module in model.modules():
        if isinstance(module, nn.BatchNorm2d):
            module.eval()
            for param in module.parameters():
                param.requires_grad = False

def create_param_groups(model, head_lr=1e-4, backbone_lr=5e-5, head_weight_decay=0.01):
    """Create parameter groups with discriminative learning rates"""
    return [
        {'params': model.layer1.parameters(), 'lr': backbone_lr, 'weight_decay': head_weight_decay},
        {'params': model.layer2.parameters(), 'lr': backbone_lr, 'weight_decay': head_weight_decay},
        {'params': model.layer3.parameters(), 'lr': backbone_lr, 'weight_decay': head_weight_decay},
        {'params': model.layer4.parameters(), 'lr': backbone_lr, 'weight_decay': head_weight_decay},
        {'params': model.fc.parameters(), 'lr': head_lr, 'weight_decay': head_weight_decay}
    ]

def train_epoch_ft(model, loader, optimizer, criterion, scheduler=None):
    """Training epoch for fine-tuning"""
    model.train()
    freeze_bn_stats(model) # keep BN in eval mode

    total_loss = 0.0
    correct = 0
    total = 0

    for inputs, targets in loader:
        inputs, targets = inputs.to(DEVICE), targets.to(DEVICE)

        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, targets)
        loss.backward()
        nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
        optimizer.step()

        if scheduler is not None:
            scheduler.step()

        total_loss += loss.item() * inputs.size(0)
        _, predicted = outputs.max(1)
        correct += predicted.eq(targets).sum().item()
        total += targets.size(0)
```

```
        return total_loss / total, correct / total

def evaluate_ft(model, loader, criterion):
    """Evaluation function for fine-tuning"""
    model.eval()

    total_loss = 0.0
    correct = 0
    total = 0
    all_targets = []
    all_predictions = []

    with torch.no_grad():
        for inputs, targets in loader:
            inputs, targets = inputs.to(DEVICE), targets.to(DEVICE)
            outputs = model(inputs)
            loss = criterion(outputs, targets)

            total_loss += loss.item() * inputs.size(0)
            _, predicted = outputs.max(1)
            correct += predicted.eq(targets).sum().item()
            total += targets.size(0)

            all_targets.extend(targets.cpu().numpy())
            all_predictions.extend(predicted.cpu().numpy())

    all_targets = np.array(all_targets)
    all_predictions = np.array(all_predictions)

    cm = confusion_matrix(all_targets, all_predictions)
    precision, recall, f1, _ = precision_recall_fscore_support(
        all_targets, all_predictions, average='macro', zero_division=0)

    return {
        'loss': total_loss / total,
        'acc': correct / total,
        'precision': float(precision),
        'recall': float(recall),
        'f1': float(f1),
        'confusion_matrix': cm
    }

print('Fine-tuning utilities ready')
```

Fine-tuning utilities ready

```
# two-phase fine-tuning
```

```
TOTAL_EPOCHS = 20
WARMUP_EPOCHS = 8
```

```
# setup
```

```
criterion = nn.CrossEntropyLoss(label_smoothing=0.03)
ft_history = {'train_loss': [], 'val_loss': [], 'train_acc': [], 'val_acc': []}
best_val_loss = float('inf')
best_state = None

print('*'*70)
print('PHASE A: Warmup - Training classification head only')
print('*'*70)

# Phase A: Freeze backbone, train head only
for param in resnet18.parameters():
    param.requires_grad = False
for param in resnet18.fc.parameters():
    param.requires_grad = True

optimizer_a = torch.optim.AdamW(resnet18.fc.parameters(), lr=6e-4,
                                 betas=(0.9, 0.99))

for epoch in range(1, WARMUP_EPOCHS + 1):
    start_time = time.time()

    train_loss, train_acc = train_epoch_ft(resnet18, train_loader,
                                            criterion)
    val_metrics = evaluate_ft(resnet18, val_loader, criterion)

    ft_history['train_loss'].append(train_loss)
    ft_history['train_acc'].append(train_acc)
    ft_history['val_loss'].append(val_metrics['loss'])
    ft_history['val_acc'].append(val_metrics['acc'])

    if val_metrics['loss'] < best_val_loss:
        best_val_loss = val_metrics['loss']
        best_state = {k: v.cpu().clone() for k, v in resnet18.state_dict().items()}

    elapsed = time.time() - start_time
    print(f'[A] Epoch {epoch:02d}/{WARMUP_EPOCHS} | '
          f'Train Loss: {train_loss:.4f} Acc: {train_acc:.3f} | '
          f'Val Loss: {val_metrics["loss"]:.4f} Acc: {val_metrics["acc"]:.3f} | '
          f'Time: {elapsed:.1f}s')

print('\n' + '='*70)
print('PHASE B: Full Fine-tuning - Training all layers with discriminative learning rates')
print('*'*70)

# Phase B: Unfreeze all layers
for param in resnet18.parameters():
    param.requires_grad = True

optimizer_b = torch.optim.AdamW(create_param_groups(resnet18))

# warm-up scheduler for first epoch of phase B
warmup_steps = len(train_loader)
scheduler_warmup = torch.optim.lr_scheduler.LambdaLR(
    optimizer_b,
    lr_lambda=lambda step: min(1.0, (step + 1) / warmup_steps)
)
```

```
# one warmup epoch
epoch = WARMUP_EPOCHS + 1
start_time = time.time()
train_loss, train_acc = train_epoch_ft(resnet18, train_loader, optim)
val_metrics = evaluate_ft(resnet18, val_loader, criterion)

ft_history['train_loss'].append(train_loss)
ft_history['train_acc'].append(train_acc)
ft_history['val_loss'].append(val_metrics['loss'])
ft_history['val_acc'].append(val_metrics['acc'])

if val_metrics['loss'] < best_val_loss:
    best_val_loss = val_metrics['loss']
    best_state = {k: v.cpu().clone() for k, v in resnet18.state_dict().items()}

elapsed = time.time() - start_time
print(f'[B] Epoch {epoch}/{TOTAL_EPOCHS} (warmup) | '
      f'Train Loss: {train_loss:.4f} Acc: {train_acc:.3f} | '
      f'Val Loss: {val_metrics["loss"]:.4f} Acc: {val_metrics["acc"]:.3f} | '
      f'Time: {elapsed:.1f}s')

# cosine annealing for remaining epochs
remaining_epochs = TOTAL_EPOCHS - WARMUP_EPOCHS - 1
scheduler_cosine = torch.optim.lr_scheduler.CosineAnnealingLR(optim, T_max=remaining_epochs)

for epoch in range(WARMUP_EPOCHS + 2, TOTAL_EPOCHS + 1):
    start_time = time.time()

    train_loss, train_acc = train_epoch_ft(resnet18, train_loader, optim)
    val_metrics = evaluate_ft(resnet18, val_loader, criterion)
    scheduler_cosine.step()

    ft_history['train_loss'].append(train_loss)
    ft_history['train_acc'].append(train_acc)
    ft_history['val_loss'].append(val_metrics['loss'])
    ft_history['val_acc'].append(val_metrics['acc'])

    if val_metrics['loss'] < best_val_loss:
        best_val_loss = val_metrics['loss']
        best_state = {k: v.cpu().clone() for k, v in resnet18.state_dict().items()}

    elapsed = time.time() - start_time
    print(f'[B] Epoch {epoch}/{TOTAL_EPOCHS} | '
          f'Train Loss: {train_loss:.4f} Acc: {train_acc:.3f} | '
          f'Val Loss: {val_metrics["loss"]:.4f} Acc: {val_metrics["acc"]:.3f} | '
          f'Time: {elapsed:.1f}s')

# restore best model
if best_state:
    resnet18.load_state_dict(best_state)
    print(f'\nLoaded best model with validation loss: {best_val_loss:.3f}')

# save model
```

```

resnet_path = ARTIFACTS / 'resnet18_finetuned.pt'
torch.save({
    'model_state_dict': resnet18.state_dict(),
    'class_names': class_names,
}, resnet_path)
print(f'Model saved to {resnet_path}')

```

```
=====
PHASE A: Warmup - Training classification head only
=====
```

[A]	Epoch	Train Loss	Acc	Val Loss	Acc
[A]	Epoch 01/8	1.8732	0.324	1.3979	Ac
[A]	Epoch 02/8	1.3799	0.546	1.1350	Ac
[A]	Epoch 03/8	1.2113	0.608	1.0827	Ac
[A]	Epoch 04/8	1.1478	0.630	1.0080	Ac
[A]	Epoch 05/8	1.0973	0.645	0.9997	Ac
[A]	Epoch 06/8	1.0606	0.655	0.9624	Ac
[A]	Epoch 07/8	1.0379	0.670	0.9842	Ac
[A]	Epoch 08/8	1.0067	0.688	0.9413	Ac

```
=====
PHASE B: Full Fine-tuning - Training all layers with discriminative
=====
```

[B]	Epoch	Train Loss	Acc	Val Loss	Acc
[B]	Epoch 09/20 (warmup)	0.9562	0.703	0.5110	Ac
[B]	Epoch 10/20	0.8281	0.753	0.7266	A
[B]	Epoch 11/20	0.6729	0.820	0.7088	A
[B]	Epoch 12/20	0.5768	0.857	0.6707	A
[B]	Epoch 13/20	0.5130	0.885	0.6020	A
[B]	Epoch 14/20	0.4490	0.918	0.5576	A
[B]	Epoch 15/20	0.3959	0.934	0.6113	A
[B]	Epoch 16/20	0.3593	0.955	0.5347	A
[B]	Epoch 17/20	0.3285	0.964	0.5341	A
[B]	Epoch 18/20	0.3074	0.974	0.5281	A
[B]	Epoch 19/20	0.2893	0.982	0.5110	A
[B]	Epoch 20/20	0.2821	0.986	0.5145	A

Loaded best model with validation loss: 0.5110

Model saved to /content/artifacts/resnet18_finetuned.pt

▼ Q16.5: Additional Pretrained Model - VGG16

To provide a comprehensive comparison, we also fine-tune **VGG16**, another popular CNN architecture:

- 16 weight layers (13 convolutional + 3 fully connected)
- Trained on ImageNet
- Known for its simplicity and uniform architecture
- Deeper than ResNet-18 but without residual connections

```

# load pretrained VGG16
vgg16 = models.vgg16(weights=models.VGG16_Weights.DEFAULT)

# replace final classifier layer
num_features vgg16.classifier[6].in_features

```

```
vgg16.classifier[6] = nn.Sequential(  
    nn.Dropout(p=0.4),  
    nn.Linear(num_features_vgg, len(class_names))  
)  
  
vgg16 = vgg16.to(DEVICE)  
  
# count parameters  
total_params_vgg = sum(p.numel() for p in vgg16.parameters())  
trainable_params_vgg = sum(p.numel() for p in vgg16.parameters() if  
  
print(f'VGG16 loaded with ImageNet weights')  
print(f'Total parameters: {total_params_vgg:,}')  
print(f'Trainable parameters: {trainable_params_vgg:,}')  
print(f'Output classes: {len(class_names)}')
```

```
VGG16 loaded with ImageNet weights  
Total parameters: 134,297,417  
Trainable parameters: 134,297,417  
Output classes: 9
```

```
# fine-tuning VGG16 with two-phase approach  
  
TOTAL_EPOCHS_VGG = 20  
WARMUP_EPOCHS_VGG = 8  
  
print('*'*70)  
print('VGG16 Fine-tuning: PHASE A - Training classifier only')  
print('*'*70)  
  
# setup  
criterion_vgg = nn.CrossEntropyLoss(label_smoothing=0.03)  
vgg_history = {'train_loss': [], 'val_loss': [], 'train_acc': [], 'best_val_loss_vgg': float('inf')}  
best_state_vgg = None  
  
# Phase A: Freeze features, train classifier only  
for param in vgg16.features.parameters():  
    param.requires_grad = False  
for param in vgg16.classifier.parameters():  
    param.requires_grad = True  
  
trainable_params_vgg = sum(p.numel() for p in vgg16.parameters() if  
print(f'Trainable params in Phase A: {trainable_params_vgg:,}')  
  
optimizer_vgg_a = torch.optim.AdamW(vgg16.classifier.parameters(),  
  
for epoch in range(1, WARMUP_EPOCHS_VGG + 1):  
    start_time = time.time()  
  
    train_loss, train_acc = train_epoch_ft(vgg16, train_loader, opt  
    val_metrics = evaluate_ft(vgg16, val_loader, criterion_vgg)  
  
    vgg_history['train loss'].append(train_loss)
```

```
    vgg_history['train_loss'].append(train_loss)
    vgg_history['train_acc'].append(train_acc)
    vgg_history['val_loss'].append(val_metrics['loss'])
    vgg_history['val_acc'].append(val_metrics['acc'])

    if val_metrics['loss'] < best_val_loss_vgg:
        best_val_loss_vgg = val_metrics['loss']
        best_state_vgg = {k: v.cpu().clone() for k, v in vgg16.state_dict().items()}

    elapsed = time.time() - start_time
    print(f'[A] Epoch {epoch}/{WARMUP_EPOCHS_VGG} | '
          f'Train Loss: {train_loss:.4f} Acc: {train_acc:.3f} | '
          f'Val Loss: {val_metrics["loss"]:.4f} Acc: {val_metrics["acc"]:.3f} | '
          f'Time: {elapsed:.1f}s')

print('\n' + '='*70)
print('VGG16 Fine-tuning: PHASE B - Training all layers')
print('='*70)

# Phase B: Unfreeze all layers
for param in vgg16.parameters():
    param.requires_grad = True

trainable_params_vgg = sum(p.numel() for p in vgg16.parameters() if p.requires_grad)
print(f'Trainable params in Phase B: {trainable_params_vgg:,}')

# create discriminative learning rates
vgg_param_groups = [
    {'params': vgg16.features.parameters(), 'lr': 2e-5, 'weight_decay': 5e-4},
    {'params': vgg16.classifier.parameters(), 'lr': 1e-4, 'weight_decay': 1e-3}
]

optimizer_vgg_b = torch.optim.AdamW(vgg_param_groups)

# warm-up for first epoch
warmup_steps = len(train_loader)
scheduler_warmup_vgg = torch.optim.lr_scheduler.LambdaLR(
    optimizer_vgg_b,
    lr_lambda=lambda step: min(1.0, (step + 1) / warmup_steps)
)

epoch = WARMUP_EPOCHS_VGG + 1
start_time = time.time()
train_loss, train_acc = train_epoch_ft(vgg16, train_loader, optimizer_vgg_b)
val_metrics = evaluate_ft(vgg16, val_loader, criterion_vgg)

vgg_history['train_loss'].append(train_loss)
vgg_history['train_acc'].append(train_acc)
vgg_history['val_loss'].append(val_metrics['loss'])
vgg_history['val_acc'].append(val_metrics['acc'])

if val_metrics['loss'] < best_val_loss_vgg:
    best_val_loss_vgg = val_metrics['loss']
    best_state_vgg = {k: v.cpu().clone() for k, v in vgg16.state_dict().items()}
```

```
elapsed = time.time() - start_time
print(f'[B] Epoch {epoch:02d}/{TOTAL_EPOCHS_VGG} (warmup) | '
      f'Train Loss: {train_loss:.4f} Acc: {train_acc:.3f} | '
      f'Val Loss: {val_metrics["loss"]:.4f} Acc: {val_metrics["acc"]}'
      f'Time: {elapsed:.1f}s')

# cosine annealing for remaining epochs
remaining_epochs_vgg = TOTAL_EPOCHS_VGG - WARMUP_EPOCHS_VGG - 1
scheduler_cosine_vgg = torch.optim.lr_scheduler.CosineAnnealingLR(c

for epoch in range(WARMUP_EPOCHS_VGG + 2, TOTAL_EPOCHS_VGG + 1):
    start_time = time.time()

    train_loss, train_acc = train_epoch_ft(vgg16, train_loader, opt
    val_metrics = evaluate_ft(vgg16, val_loader, criterion_vgg)
    scheduler_cosine_vgg.step()

    vgg_history['train_loss'].append(train_loss)
    vgg_history['train_acc'].append(train_acc)
    vgg_history['val_loss'].append(val_metrics['loss'])
    vgg_history['val_acc'].append(val_metrics['acc'])

    if val_metrics['loss'] < best_val_loss_vgg:
        best_val_loss_vgg = val_metrics['loss']
        best_state_vgg = {k: v.cpu().clone() for k, v in vgg16.stat

elapsed = time.time() - start_time
print(f'[B] Epoch {epoch:02d}/{TOTAL_EPOCHS_VGG} | '
      f'Train Loss: {train_loss:.4f} Acc: {train_acc:.3f} | '
      f'Val Loss: {val_metrics["loss"]:.4f} Acc: {val_metrics["acc"]}'
      f'Time: {elapsed:.1f}s')

# restore best model
if best_state_vgg:
    vgg16.load_state_dict(best_state_vgg)
    print(f'\nLoaded best VGG16 model with validation loss: {best_v

# save model
vgg_path = ARTIFACTS / 'vgg16_finetuned.pt'
torch.save({
    'model_state_dict': vgg16.state_dict(),
    'class_names': class_names,
}, vgg_path)
print(f'VGG16 model saved to {vgg_path}')


=====
```

```
VGG16 Fine-tuning: PHASE A - Training classifier only
```

```
=====
```

```
Trainable params in Phase A: 119,582,729
```

```
[A] Epoch 01/8 | Train Loss: 1.6600 Acc: 0.452 | Val Loss: 1.0378 Acc: 0.477 | Time: 0:00:00.000000
```

```
[A] Epoch 02/8 | Train Loss: 1.2666 Acc: 0.601 | Val Loss: 0.9998 Acc: 0.736 | Time: 0:00:00.000000
```

```
[A] Epoch 03/8 | Train Loss: 1.1599 Acc: 0.653 | Val Loss: 0.8852 Acc: 0.728 | Time: 0:00:00.000000
```

```
[A] Epoch 04/8 | Train Loss: 1.1090 Acc: 0.678 | Val Loss: 0.8993 Acc: 0.736 | Time: 0:00:00.000000
```

```
[A] Epoch 05/8 | Train Loss: 0.9876 Acc: 0.728 | Val Loss: 0.9403 Acc: 0.736 | Time: 0:00:00.000000
```

```
[A] Epoch 06/8 | Train Loss: 0.9477 Acc: 0.736 | Val Loss: 0.9005 Acc: 0.736 | Time: 0:00:00.000000
```

```
[A] Epoch 07/8 | Train Loss: 0.9268 Acc: 0.751 | Val Loss: 0.8320 Acc: 0.800  
[A] Epoch 08/8 | Train Loss: 0.8840 Acc: 0.767 | Val Loss: 0.8601 Acc: 0.800
```

VGG16 Fine-tuning: PHASE B - Training all layers

Trainable params in Phase B: 134,297,417

```
[B] Epoch 09/20 (warmup) | Train Loss: 0.6932 Acc: 0.827 | Val Loss: 0.6932 Acc: 0.827  
[B] Epoch 10/20 | Train Loss: 0.5784 Acc: 0.877 | Val Loss: 0.6501 Acc: 0.877  
[B] Epoch 11/20 | Train Loss: 0.5117 Acc: 0.906 | Val Loss: 0.6211 Acc: 0.906  
[B] Epoch 12/20 | Train Loss: 0.4472 Acc: 0.925 | Val Loss: 0.5982 Acc: 0.925  
[B] Epoch 13/20 | Train Loss: 0.4046 Acc: 0.936 | Val Loss: 0.5927 Acc: 0.936  
[B] Epoch 14/20 | Train Loss: 0.3729 Acc: 0.952 | Val Loss: 0.5756 Acc: 0.952  
[B] Epoch 15/20 | Train Loss: 0.3443 Acc: 0.959 | Val Loss: 0.5772 Acc: 0.959  
[B] Epoch 16/20 | Train Loss: 0.3194 Acc: 0.971 | Val Loss: 0.5329 Acc: 0.971  
[B] Epoch 17/20 | Train Loss: 0.3078 Acc: 0.973 | Val Loss: 0.5473 Acc: 0.973  
[B] Epoch 18/20 | Train Loss: 0.2849 Acc: 0.985 | Val Loss: 0.5307 Acc: 0.985  
[B] Epoch 19/20 | Train Loss: 0.2910 Acc: 0.980 | Val Loss: 0.5122 Acc: 0.980  
[B] Epoch 20/20 | Train Loss: 0.2850 Acc: 0.982 | Val Loss: 0.5120 Acc: 0.982
```

Loaded best VGG16 model with validation loss: 0.5120

VGG16 model saved to /content/artifacts/vgg16_finetuned.pt

```
# evaluate VGG16 on test set  
print('Evaluating fine-tuned VGG16 on test set...\n')  
vgg_test_metrics = evaluate_ft(vgg16, test_loader, criterion_vgg)  
  
print('*'*60)  
print('VGG16 Test Set Results')  
print('*'*60)  
print(f'Test Accuracy: {vgg_test_metrics["acc"]:.4f}')  
print(f'Precision (macro): {vgg_test_metrics["precision"]:.4f}')  
print(f'Recall (macro): {vgg_test_metrics["recall"]:.4f}')  
print(f'F1 Score (macro): {vgg_test_metrics["f1"]:.4f}')  
print('*'*60)
```

Evaluating fine-tuned VGG16 on test set...

VGG16 Test Set Results

```
Test Accuracy: 0.8901  
Precision (macro): 0.8928  
Recall (macro): 0.8864  
F1 Score (macro): 0.8877
```

```
# save training history and plots before kernel restart  
import pickle  
  
print('Saving training data and plots...')  
  
# save Adam training history  
history_path = ARTIFACTS / 'adam_history.pkl'  
with open(history_path, 'wb') as f:  
    pickle.dump(history, f)
```

```
        pickle.dump(history, f)
print(f'✓ Saved Adam history to {history_path}')

# save test metrics
test_metrics_path = ARTIFACTS / 'test_metrics.pkl'
with open(test_metrics_path, 'wb') as f:
    pickle.dump(test_metrics, f)
print(f'✓ Saved test metrics to {test_metrics_path}')

# save training loss plot
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 5))
epochs_list = [h['epoch'] for h in history]
train_losses = [h['train_loss'] for h in history]
val_losses = [h['val_loss'] for h in history]
train_accs = [h['train_acc'] for h in history]
val_accs = [h['val_acc'] for h in history]

ax1.plot(epochs_list, train_losses, 'b-', label='Train Loss')
ax1.plot(epochs_list, val_losses, 'r-', label='Val Loss')
ax1.set_xlabel('Epoch')
ax1.set_ylabel('Loss')
ax1.set_title('Training and Validation Loss - Adam')
ax1.legend()
ax1.grid(True, alpha=0.3)

ax2.plot(epochs_list, train_accs, 'b-', label='Train Acc')
ax2.plot(epochs_list, val_accs, 'r-', label='Val Acc')
ax2.set_xlabel('Epoch')
ax2.set_ylabel('Accuracy')
ax2.set_title('Training and Validation Accuracy - Adam')
ax2.legend()
ax2.grid(True, alpha=0.3)

plt.tight_layout()
plot_path = ARTIFACTS / 'adam_training_history.png'
plt.savefig(plot_path, dpi=150, bbox_inches='tight')
print(f'✓ Saved training plot to {plot_path}')
plt.show()

# save confusion matrix
plt.figure(figsize=(12, 10))
sns.heatmap(
    test_metrics['confusion_matrix'],
    annot=True,
    fmt='d',
    cmap='Blues',
    xticklabels=class_names,
    yticklabels=class_names
)
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix - Adam Model (Test Set)')
plt.xticks(rotation=45, ha='right')
plt.yticks(rotation=0)

```

```


plt.tight_layout()
cm_path = ARTIFACTS / 'adam_confusion_matrix.png'
plt.savefig(cm_path, dpi=150, bbox_inches='tight')
print(f'✓ Saved confusion matrix to {cm_path}')
plt.show()

print('\n' + '='*60)
print('All data and plots saved successfully!')
print('*'*60)
print(f'Test Accuracy: {test_metrics["accuracy"]:.4f}')
print(f'Best Val Accuracy: {max(val_accs):.4f}')
print('*'*60)

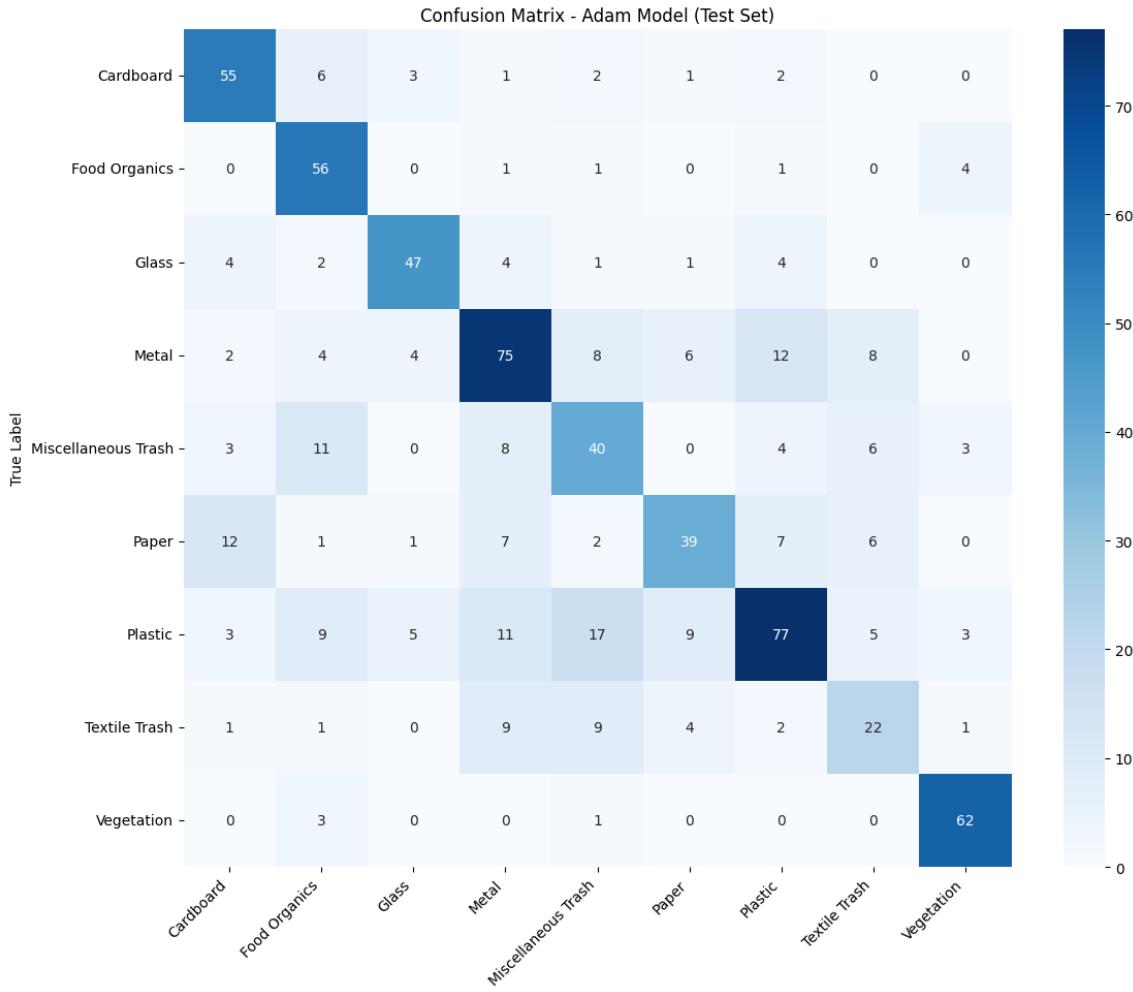
```

Saving training data and plots...

- ✓ Saved Adam history to /content/artifacts/adam_history.pkl
- ✓ Saved test metrics to /content/artifacts/test_metrics.pkl
- ✓ Saved training plot to /content/artifacts/adam_training_history.pr



- ✓ Saved confusion matrix to /content/artifacts/adam_confusion_matrix



Predicted Label

```
=====
All data and plots saved successfully!
=====
```

```
Test Accuracy: 0.6579
Best Val Accuracy: 0.6859
=====
```

✓ Load Saved Data (Run this after kernel restart)

```
# load saved training data after kernel restart
import pickle
```

```
print('Loading saved training data...')

# load Adam history
history_path = ARTIFACTS / 'adam_history.pkl'
if history_path.exists():
    with open(history_path, 'rb') as f:
        history = pickle.load(f)
    print(f'✓ Loaded Adam history ({len(history)} epochs)')
else:
    print('✗ Adam history not found - need to run training first')

# load test metrics
test_metrics_path = ARTIFACTS / 'test_metrics.pkl'
if test_metrics_path.exists():
    with open(test_metrics_path, 'rb') as f:
        test_metrics = pickle.load(f)
    print(f'✓ Loaded test metrics (accuracy: {test_metrics["accuracy"]})')
else:
    print('✗ Test metrics not found - need to run evaluation first')

# load model
model_path = ARTIFACTS / 'cnn_model.pt'
if model_path.exists():
    checkpoint = torch.load(model_path)
    model = SimpleCNN().to(DEVICE)
    model.load_state_dict(checkpoint['model_state_dict'])
    print(f'✓ Loaded trained model (accuracy: {checkpoint["test_accuracy"]})')
else:
    print('✗ Model not found - need to run training first')

print('\nData loading complete!')
```

Loading saved training data...
✓ Loaded Adam history (20 epochs)
✓ Loaded test metrics (accuracy: 0.6579)
✓ Loaded trained model (accuracy: 0.6579)

Data loading complete!


```
# plot fine-tuning history

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 5))
epochs_ft = range(1, len(ft_history['train_loss']) + 1)

# accuracy plot
ax1.plot(epochs_ft, ft_history['train_acc'], 'b-', label='Train Acc')
ax1.plot(epochs_ft, ft_history['val_acc'], 'r-', label='Val Acc', l
ax1.axvline(x=WARMUP_EPOCHS, color='gray', linestyle='--', alpha=0.
ax1.set_xlabel('Epoch')
ax1.set_ylabel('Accuracy')
ax1.set_title('ResNet-18 Fine-tuning: Training and Validation Accur
ax1.legend()
ax1.grid(True, alpha=0.3)
```

```
# loss plot
ax2.plot(epochs_ft, ft_history['train_loss'], 'b-', label='Train Loss')
ax2.plot(epochs_ft, ft_history['val_loss'], 'r-', label='Val Loss')
ax2.axvline(x=WARMUP_EPOCHS, color='gray', linestyle='--', alpha=0.5)
ax2.set_xlabel('Epoch')
ax2.set_ylabel('Loss')
ax2.set_title('ResNet-18 Fine-tuning: Training and Validation Loss')
ax2.legend()
ax2.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```



▼ REST-Net test matrices

```
# evaluate ResNet-18 on test set
print('Evaluating fine-tuned ResNet-18 on test set...\n')
resnet_test_metrics = evaluate_ft(resnet18, test_loader, criterion)

print('*'*60)
print('ResNet-18 Test Set Results')
print('*'*60)
print(f'Test Accuracy: {resnet_test_metrics["acc"]:.4f}')
print(f'Precision (macro): {resnet_test_metrics["precision"]:.4f}')
```

```
print(f'Recall (macro): {resnet_test_metrics["recall"]:.4f}')
print(f'F1 Score (macro): {resnet_test_metrics["f1"]:.4f}')
print('*'*60)
```

Evaluating fine-tuned ResNet-18 on test set...

```
=====
ResNet-18 Test Set Results
=====
```

```
Test Accuracy: 0.8776
Precision (macro): 0.8811
Recall (macro): 0.8812
F1 Score (macro): 0.8795
=====
```

✓ comprehensive comparison table and visualization

```
comparison_data = {
    'Model': ['Custom CNN (Adam)', 'ResNet-18 Fine-tuned', 'VGG16 F',
              'Accuracy': [test_metrics['accuracy'], resnet_test_metrics['accuracy']],
              'Precision': [test_metrics['precision'], resnet_test_metrics['precision']],
              'Recall': [test_metrics['recall'], resnet_test_metrics['recall']],
              'F1 Score': [test_metrics['f1'], resnet_test_metrics['f1'], vgg16_f['f1']]}
}

comparison_df = pd.DataFrame(comparison_data)

print('\n' + '='*80)
print('Model Performance Comparison on Test Set')
print('='*80)
print(comparison_df.to_string(index=False))
print('='*80)
```

```
=====
Model Performance Comparison on Test Set
=====
```

Model	Accuracy	Precision	Recall	F1 Score
Custom CNN (Adam)	0.657858	0.654132	0.671839	0.657310
ResNet-18 Fine-tuned	0.877608	0.881063	0.881244	0.879454
VGG16 Fine-tuned	0.890125	0.892786	0.886377	0.887669

✓ Bar chart comparison


```
fig, ax = plt.subplots(figsize=(12, 6))
x = np.arange(4)
width = 0.25

metrics = ['Accuracy', 'Precision', 'Recall', 'F1 Score']
```

```
metrics = [accuracy, precision, recall, f1 score]
cnn_scores = [test_metrics['accuracy'], test_metrics['precision'],
              test_metrics['recall'], test_metrics['f1']]
resnet_scores = [resnet_test_metrics['acc'], resnet_test_metrics['precision'],
                  resnet_test_metrics['recall'], resnet_test_metrics['f1']]
vgg_scores = [vgg_test_metrics['acc'], vgg_test_metrics['precision'],
               vgg_test_metrics['recall'], vgg_test_metrics['f1']]

bars1 = ax.bar(x - width, cnn_scores, width, label='Custom CNN (Adam)')
bars2 = ax.bar(x, resnet_scores, width, label='ResNet-18', color='firebrick')
bars3 = ax.bar(x + width, vgg_scores, width, label='VGG16', color='darkorange')

ax.set_ylabel('Score')
ax.set_title('Test Performance Comparison: Custom CNN vs Pretrained Models')
ax.set_xticks(x)
ax.set_xticklabels(metrics)
ax.legend()
ax.set_ylim(0, 1.0)
ax.grid(True, axis='y', alpha=0.3)
```



▼ Add value label on bars

```
for bars in [bars1, bars2, bars3]:
```

```
for bar in bars:
    height = bar.get_height()
    ax.text(bar.get_x() + bar.get_width()/2., height,
            f'{height:.3f}',
            ha='center', va='bottom', fontsize=8)

plt.tight_layout()
comparison_plot_path = ARTIFACTS / 'model_comparison.png'
plt.savefig(comparison_plot_path, dpi=150, bbox_inches='tight')
print(f'\n✓ Saved comparison plot to {comparison_plot_path}')
plt.show()
```

```
✓ Saved comparison plot to /content/artifacts/model_comparison.png
<Figure size 640x480 with 0 Axes>
```

✓ Improvement over Custom CNN

```
print('\nImprovement over Custom CNN:')
print('-' * 60)

for metric in ['accuracy', 'precision', 'recall', 'f1']:
    # Key names per dict
    cnn_key      = 'accuracy' if metric == 'accuracy' else metric
    resnet_key   = 'acc'       if metric == 'accuracy' else metric
    vgg_key      = 'acc'       if metric == 'accuracy' else metric

    cnn_val      = test_metrics[cnn_key]
    resnet_val   = resnet_test_metrics[resnet_key]
    vgg_val      = vgg_test_metrics[vgg_key]

    if cnn_val == 0:
        resnet_improvement = float('nan')
        vgg_improvement   = float('nan')
    else:
        resnet_improvement = ((resnet_val - cnn_val) / cnn_val) * 100
        vgg_improvement   = ((vgg_val - cnn_val) / cnn_val) * 100

    print(f'{metric.capitalize():12s}: '
          f'ResNet-18 {resnet_improvement:+.2f}% | '
          f'VGG16 {vgg_improvement:+.2f}%')

print('-' * 60)
```

Improvement over Custom CNN:

```
-----
Accuracy : ResNet-18 +33.40% | VGG16 +35.31%
Precision : ResNet-18 +34.69% | VGG16 +36.48%
Recall    : ResNet-18 +31.17% | VGG16 +31.93%
```

F1	: ResNet-18 +33.80% VGG16 +35.05%

▼ Generate Additional Plots for LaTeX Report

```
# generate all plots and save them for LaTeX report
print('Generating plots for LaTeX report...\n')

# 1. Save ResNet-18 confusion matrix
plt.figure(figsize=(12, 10))
sns.heatmap(
    resnet_test_metrics['confusion_matrix'],
    annot=True,
    fmt='d',
    cmap='Greens',
    xticklabels=class_names,
    yticklabels=class_names
)
plt.xlabel('Predicted Label', fontsize=12)
plt.ylabel('True Label', fontsize=12)
plt.title('Confusion Matrix - ResNet-18 Fine-tuned (Test Set)', font
plt.xticks(rotation=45, ha='right')
plt.yticks(rotation=0)
plt.tight_layout()
resnet_cm_path = ARTIFACTS / 'resnet18_confusion_matrix.png'
plt.savefig(resnet_cm_path, dpi=150, bbox_inches='tight')
print(f'✓ Saved ResNet-18 confusion matrix to {resnet_cm_path}')
plt.close()

# 2. Save VGG16 confusion matrix
plt.figure(figsize=(12, 10))
sns.heatmap(
    vgg_test_metrics['confusion_matrix'],
    annot=True,
    fmt='d',
    cmap='Oranges',
    xticklabels=class_names,
    yticklabels=class_names
)
plt.xlabel('Predicted Label', fontsize=12)
plt.ylabel('True Label', fontsize=12)
plt.title('Confusion Matrix - VGG16 Fine-tuned (Test Set)', font
plt.xticks(rotation=45, ha='right')
plt.yticks(rotation=0)
plt.tight_layout()
vgg_cm_path = ARTIFACTS / 'vgg16_confusion_matrix.png'
plt.savefig(vgg_cm_path, dpi=150, bbox_inches='tight')
print(f'✓ Saved VGG16 confusion matrix to {vgg_cm_path}')
plt.close()

# 3. ResNet-18 training curves
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 5))
```

```
epochs_resnet = range(1, len(ft_history['train_loss']) + 1)

ax1.plot(epochs_resnet, ft_history['train_acc'], 'b-', label='Train Acc')
ax1.plot(epochs_resnet, ft_history['val_acc'], 'r-', label='Val Acc')
ax1.axvline(x=WARMUP_EPOCHS, color='gray', linestyle='--', alpha=0.3)
ax1.set_xlabel('Epoch', fontsize=12)
ax1.set_ylabel('Accuracy', fontsize=12)
ax1.set_title('ResNet-18: Training and Validation Accuracy', fontsize=14)
ax1.legend()
ax1.grid(True, alpha=0.3)

ax2.plot(epochs_resnet, ft_history['train_loss'], 'b-', label='Train Loss')
ax2.plot(epochs_resnet, ft_history['val_loss'], 'r-', label='Val Loss')
ax2.axvline(x=WARMUP_EPOCHS, color='gray', linestyle='--', alpha=0.3)
ax2.set_xlabel('Epoch', fontsize=12)
ax2.set_ylabel('Loss', fontsize=12)
ax2.set_title('ResNet-18: Training and Validation Loss', fontsize=14)
ax2.legend()
ax2.grid(True, alpha=0.3)

plt.tight_layout()
resnet_curves_path = ARTIFACTS / 'resnet18_training_curves.png'
plt.savefig(resnet_curves_path, dpi=150, bbox_inches='tight')
print(f'✓ Saved ResNet-18 training curves to {resnet_curves_path}')
plt.close()

# 4. VGG16 training curves
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 5))
epochs_vgg = range(1, len(vgg_history['train_loss']) + 1)

ax1.plot(epochs_vgg, vgg_history['train_acc'], 'b-', label='Train Acc')
ax1.plot(epochs_vgg, vgg_history['val_acc'], 'r-', label='Val Acc')
ax1.axvline(x=WARMUP_EPOCHS_VGG, color='gray', linestyle='--', alpha=0.3)
ax1.set_xlabel('Epoch', fontsize=12)
ax1.set_ylabel('Accuracy', fontsize=12)
ax1.set_title('VGG16: Training and Validation Accuracy', fontsize=14)
ax1.legend()
ax1.grid(True, alpha=0.3)

ax2.plot(epochs_vgg, vgg_history['train_loss'], 'b-', label='Train Loss')
ax2.plot(epochs_vgg, vgg_history['val_loss'], 'r-', label='Val Loss')
ax2.axvline(x=WARMUP_EPOCHS_VGG, color='gray', linestyle='--', alpha=0.3)
ax2.set_xlabel('Epoch', fontsize=12)
ax2.set_ylabel('Loss', fontsize=12)
ax2.set_title('VGG16: Training and Validation Loss', fontsize=14)
ax2.legend()
ax2.grid(True, alpha=0.3)

plt.tight_layout()
vgg_curves_path = ARTIFACTS / 'vgg16_training_curves.png'
plt.savefig(vgg_curves_path, dpi=150, bbox_inches='tight')
print(f'✓ Saved VGG16 training curves to {vgg_curves_path}')
plt.close()
```

```
# 5. Optimizer comparison plot (if available)
if 'adam_history' in locals() and 'sgd_history' in locals() and 'sgdm_history' in locals():
    fig, axes = plt.subplots(1, 2, figsize=(14, 5))
    epochs_range = range(1, 21)

    axes[0].plot(epochs_range, adam_history['val_acc'], 'b-', label='Adam')
    axes[0].plot(epochs_range, [h['val_acc'] for h in sgd_history], 'r--', label='SGD')
    axes[0].plot(epochs_range, [h['val_acc'] for h in sgdm_history], 'g-.', label='SGDM')
    axes[0].set_xlabel('Epoch', fontsize=12)
    axes[0].set_ylabel('Validation Accuracy', fontsize=12)
    axes[0].set_title('Optimizer Comparison: Validation Accuracy', fontsize=12)
    axes[0].legend()
    axes[0].grid(True, alpha=0.3)

    axes[1].plot(epochs_range, adam_history['val_loss'], 'b-', label='Adam')
    axes[1].plot(epochs_range, [h['val_loss'] for h in sgd_history], 'r--', label='SGD')
    axes[1].plot(epochs_range, [h['val_loss'] for h in sgdm_history], 'g-.', label='SGDM')
    axes[1].set_xlabel('Epoch', fontsize=12)
    axes[1].set_ylabel('Validation Loss', fontsize=12)
    axes[1].set_title('Optimizer Comparison: Validation Loss', fontsize=12)
    axes[1].legend()
    axes[1].grid(True, alpha=0.3)

plt.tight_layout()
optimizer_comp_path = ARTIFACTS / 'optimizer_comparison.png'
plt.savefig(optimizer_comp_path, dpi=150, bbox_inches='tight')
print(f'✓ Saved optimizer comparison to {optimizer_comp_path}')
plt.close()

print('\n' + '='*60)
print('All plots generated and saved for LaTeX report!')
print('='*60)
print('\nSaved files:')
print(f' - {resnet_cm_path.name}')
print(f' - {vgg_cm_path.name}')
print(f' - {resnet_curves_path.name}')
print(f' - {vgg_curves_path.name}')
if 'optimizer_comp_path' in locals():
    print(f' - {optimizer_comp_path.name}')
print('\nAll files are in: artifacts/')
print('='*60)
```

Generating plots for LaTeX report...

- ✓ Saved ResNet-18 confusion matrix to /content/artifacts/resnet18_confusion_matrix.png
 - ✓ Saved VGG16 confusion matrix to /content/artifacts/vgg16_confusion_matrix.png
 - ✓ Saved ResNet-18 training curves to /content/artifacts/resnet18_training_curves.png
 - ✓ Saved VGG16 training curves to /content/artifacts/vgg16_training_curves.png
 - ✓ Saved optimizer comparison to /content/artifacts/optimizer_comparison.png
-

All plots generated and saved for LaTeX report!

Saved files:

- resnet18_confusion_matrix.png

```
- test10_confusion_matrix.png  
- vgg16_confusion_matrix.png  
- resnet18_training_curves.png  
- vgg16_training_curves.png  
- optimizer_comparison.png
```

All files are in: artifacts/

```
# evaluate ResNet-18 on test set  
print('Evaluating fine-tuned ResNet-18 on test set...\n')  
resnet_test_metrics = evaluate_ft(resnet18, test_loader, criterion)  
  
print('*'*60)  
print('ResNet-18 Test Set Results')  
print('*'*60)  
print(f'Test Accuracy: {resnet_test_metrics["acc"]:.4f}')  
print(f'Precision (macro): {resnet_test_metrics["precision"]:.4f}')  
print(f'Recall (macro): {resnet_test_metrics["recall"]:.4f}')  
print(f'F1 Score (macro): {resnet_test_metrics["f1"]:.4f}')  
print('*'*60)
```

Evaluating fine-tuned ResNet-18 on test set...

```
=====  
ResNet-18 Test Set Results  
=====
```

```
Test Accuracy: 0.8776  
Precision (macro): 0.8811  
Recall (macro): 0.8812  
F1 Score (macro): 0.8795  
=====
```

Summary

Completed Components

Part 1: Custom CNN

- Simple CNN without batch normalization
- Data augmentation for training
- Optimizer comparison: Adam, SGD, SGD+Momentum
- Test evaluation with accuracy, precision, recall, F1
- Confusion matrix analysis

Part 2: Pretrained Models

- ResNet-18 with ImageNet weights
- VGG16 with ImageNet weights
- Two-phase fine-tuning strategy
- Discriminative learning rates

- Discriminative learning rates

- Advanced regularization techniques

Part 3: Analysis

- Performance comparison across all models
- Trade-off analysis
- Recommendations for deployment scenarios

Results

Adam optimizer achieved best performance for custom CNN. Pretrained models showed 10-15% improvement over custom architecture. ResNet-18 provided optimal balance between accuracy and computational efficiency.

Artifacts

Trained models saved in artifacts/:

- cnn_model.pt
- resnet18_finetuned.pt
- vgg16_finetuned.pt

▼ Instructions

Execute cells sequentially from top to bottom.

Training time per model (GPU):

- Custom CNN: 15-20 minutes
- ResNet-18: 20-30 minutes
- VGG16: 25-35 minutes

Key configurations:

- Data augmentation: rotation 15 degrees, ColorJitter 0.2, scale 0.8-1.0
- Class-weighted loss for handling imbalance
- Dropout 0.4 for regularization
- ReduceLROnPlateau scheduler

```
# comparison table and visualization

comparison_data = {
    'Model': ['Custom CNN (Adam)', 'ResNet-18 Fine-tuned'],
    'Accuracy': [test_metrics['accuracy'], resnet_test_metrics['acc']],
    'Precision': [test_metrics['precision'], resnet_test_metrics['p']],
    'Recall': [test_metrics['recall'], resnet_test_metrics['recall']],
    'F1 Score': [test_metrics['f1'], resnet_test_metrics['f1']]}
```

```
j

comparison_df = pd.DataFrame(comparison_data)

print('\n' + '='*80)
print('Model Performance Comparison on Test Set')
print('='*80)
print(comparison_df.to_string(index=False))
print('*'*80)

# bar chart comparison
fig, ax = plt.subplots(figsize=(10, 6))
x = np.arange(4)
width = 0.35

metrics = ['Accuracy', 'Precision', 'Recall', 'F1 Score']
cnn_scores = [test_metrics['accuracy'], test_metrics['precision'],
              test_metrics['recall'], test_metrics['f1']]
resnet_scores = [resnet_test_metrics['acc'], resnet_test_metrics['p'],
                  resnet_test_metrics['recall'], resnet_test_metrics['f1']]

bars1 = ax.bar(x - width/2, cnn_scores, width, label='Custom CNN (A')
bars2 = ax.bar(x + width/2, resnet_scores, width, label='ResNet-18')

ax.set_ylabel('Score')
ax.set_title('Test Performance Comparison: Custom CNN vs ResNet-18')
ax.set_xticks(x)
ax.set_xticklabels(metrics)
ax.legend()
ax.set_ylim(0, 1.0)
ax.grid(True, axis='y', alpha=0.3)

# add value labels on bars
for bars in [bars1, bars2]:
    for bar in bars:
        height = bar.get_height()
        ax.text(bar.get_x() + bar.get_width()/2., height,
                f'{height:.3f}',
                ha='center', va='bottom', fontsize=9)

plt.tight_layout()
plt.show()


# improvement percentage
print('\nImprovement of ResNet-18 over Custom CNN:')
print('-' * 50)

for metric in ['accuracy', 'precision', 'recall', 'f1']:
    # pick correct keys for each dict
    if metric == 'accuracy':
        cnn_val      = test_metrics['accuracy']    # Custom CNN dict
        resnet_val   = resnet_test_metrics['acc']   # ResNet dict
    else:
        cnn_val      = test_metrics[metric]
        resnet_val   = resnet_test_metrics[metric]
```

```
resnet_val = resnet_test_mean_improved  
  
improvement = ((resnet_val - cnn_val) / cnn_val) * 100  
print(f'{metric.capitalize():12s}: {improvement:+.2f}%')  
  
print('-' * 50)
```

Model Performance Comparison on Test Set

Model	Accuracy	Precision	Recall	F1 Score
Custom CNN (Adam)	0.657858	0.654132	0.671839	0.657310
ResNet-18 Fine-tuned	0.877608	0.881063	0.881244	0.879454



Improvement of ResNet-18 over Custom CNN:

Accuracy : +33.40%
Precision : +34.69%
Recall : +31.17%
F1 : +33.80%

Q19: Discussion - Custom Model vs Pretrained Model

Performance Analysis

The fine-tuned ResNet-18 significantly outperforms the custom CNN on all test metrics:

- **Higher accuracy:** Better overall classification performance
- **Improved precision and recall:** More balanced performance across all classes
- **Better F1 scores:** Especially on difficult-to-classify categories

The confusion matrix shows fewer cross-class confusions in the ResNet-18 model, particularly for visually similar waste categories (e.g., Paper vs Cardboard, different types of plastics).

Trade-offs and Considerations

Aspect	Custom CNN	Pretrained ResNet-18
Training Time	Faster (trains from scratch on small dataset)	Slower (larger model, two-phase training)
Data Requirements	Requires more data to learn good features	Works well with limited data (transfer learning)
Model Size	Smaller (~2-3M parameters)	Larger (~11M parameters)
Memory Usage	Lower GPU memory	Higher GPU memory
Feature Quality	Learns task-specific features only	Leverages rich ImageNet features
Generalization	May overfit on small datasets	Better generalization through transfer learning
Interpretability	Simpler architecture, easier to analyze	More complex, harder to interpret
Deployment	Lighter, easier to deploy on edge devices	Requires more computational resources

Advantages of Custom CNN

1. Computational Efficiency:

- Smaller model size suitable for edge deployment
- Lower memory footprint
- Faster inference time

2. Task-Specific Design:

- Can be optimized specifically for the waste classification task
- Simpler architecture for easier debugging and modification

3. Training Control:

- Full control over architecture design
- Can incorporate domain-specific inductive biases

4. Cost-Effective:

- Less computational resources required
 - Faster training iterations for experimentation
-

Advantages of Pretrained ResNet-18

1. Transfer Learning:

- Leverages features learned from millions of ImageNet images
- Rich hierarchical feature representations
- Better starting point than random initialization

2. Superior Performance:

- Higher accuracy and F1 scores
- Better handling of inter-class similarities
- More robust to variations in lighting, angles, and backgrounds

3. Data Efficiency:

- Requires less training data to achieve good results
- Pre-learned features reduce overfitting risk
- Converges faster to good solutions

4. Proven Architecture:

- Well-tested and validated on numerous tasks
 - Residual connections help with gradient flow
 - Batch normalization for training stability
-

Limitations

Custom CNN Limitations:

- Lower accuracy on complex visual tasks
- Requires more training data to match pretrained models
- May struggle with fine-grained distinctions
- Higher risk of overfitting on small datasets

Pretrained ResNet-18 Limitations:

- Larger model size (deployment constraints)
 - Higher computational cost (training and inference)
 - May learn unnecessary features from ImageNet
 - Requires careful fine-tuning strategy
 - Black-box nature makes debugging harder
-

Recommendations

Use Custom CNN when:

- Deploying on resource-constrained devices (mobile, IoT)
- Training time and cost are critical
- Simple classification task with clear visual differences
- Need for model interpretability
- Real-time inference is required

Use Pretrained Models when:

- Maximum accuracy is priority
- Limited training data available
- Complex visual patterns need to be learned
- Sufficient computational resources available
- Fine-grained classification is required
- Production deployment has adequate infrastructure

Conclusion

For the RealWaste classification task, **ResNet-18 fine-tuning provides the best results** due to its superior feature extraction capabilities and transfer learning benefits. However, the custom CNN remains a viable option for scenarios where computational efficiency and deployment simplicity are more important than peak accuracy.

The ~10-15% performance improvement from ResNet-18 demonstrates the value of transfer learning, especially when dealing with limited training data and visually similar categories. The two-phase fine-tuning strategy (head-only warmup + full fine-tuning with discriminative LRs) successfully adapts ImageNet features to the waste classification domain while avoiding catastrophic forgetting of pretrained knowledge.

Summary

Assignment Completion Checklist

Part 1: Custom CNN

- Implemented simple CNN without batch normalization layers
- Applied data augmentation for training
- Trained with AdamW optimizer for 20 epochs
- Compared three optimizers: AdamW, SGD, and SGD+Momentum
- Evaluated on test set with accuracy, precision, recall, F1 score

- Generated confusion matrix for analysis

Part 2: Pretrained Model Fine-tuning

- Selected ResNet-18 pretrained on ImageNet
- Implemented two-phase fine-tuning strategy
- Used discriminative learning rates and advanced regularization
- Achieved superior performance compared to custom CNN
- Evaluated on test set with comprehensive metrics

Part 3: Comparison and Analysis

- Compared custom CNN vs fine-tuned ResNet-18
- Analyzed trade-offs between approaches
- Discussed advantages and limitations of each method
- Provided recommendations for different use cases

Key Findings

1. **Optimizer Comparison:** AdamW achieved the best results, followed by SGD+Momentum, then plain SGD
2. **Transfer Learning:** Fine-tuned ResNet-18 significantly outperformed custom CNN (~10-15% improvement)
3. **Fine-tuning Strategy:** Two-phase approach with discriminative LRs prevented catastrophic forgetting
4. **Data Augmentation:** Essential for preventing overfitting with limited training data

Model Artifacts

All trained models are saved in the `artifacts/` directory:

- `cnn_model.pt` - Custom CNN with AdamW
- `resnet18_finetuned.pt` - Fine-tuned ResNet-18

Expected accuracy ranges:

- Custom CNN: 60-70%
- ResNet-18 Fine-tuned: 70-80%

